

Scenario 2: Real-time Network Monitoring using **Azure Cloud**, we'll design an end-to-end streaming data pipeline leveraging Azure services such as **Azure Event Hubs**, **Azure Databricks**, **Azure Data Lake Storage (ADLS) Gen2**, **Azure Functions**, and **Power BI/Fabric**. This solution will enable you to collect, process, analyze, and visualize network performance metrics in real-time, as well as trigger alerts for anomalies and performance degradation.

Overview of the Solution

1. **Data Ingestion:** Stream network performance data from devices using **Azure Event Hubs**.
 2. **Stream Processing and Enrichment:** Use **Azure Databricks** with Structured Streaming to process and enrich the data.
 3. **Storage:** Store processed data in **ADLS Gen2**.
 4. **Real-time Alerting:** Implement anomaly detection and trigger alerts using **Azure Databricks** and **Azure Functions**.
 5. **Visualization:** Create real-time dashboards and reports using **Power BI/Fabric**.
 6. **Monitoring and Automation:** Monitor the pipeline and automate processes using **Azure Monitor** and **Azure Data Factory (ADF)**.
-

Step-by-Step Implementation

1. Data Ingestion with Azure Event Hubs

Objective: Stream network performance data from various devices (e.g., routers) in real-time.

Steps:

1. **Create an Azure Event Hubs Namespace and Event Hub:**
 - **Navigate to Azure Portal:**
 - Search for **Event Hubs** and create a new **Event Hubs Namespace**.
 - Within the namespace, create an **Event Hub** (e.g., network-performance-hub).
 - **Configure Event Hub:**
 - Set the **Partition Count** based on expected throughput.

- Configure **Capture** settings if you want to automatically capture streaming data to ADLS Gen2.

2. Configure Devices to Send Data to Event Hub:

- **Device Configuration:**

- Ensure that your network devices (routers, etc.) can send data to Azure Event Hubs. This typically involves setting up an **IoT Edge** or using **SDKs/APIs** to publish events.

- **Data Format:**

- Ensure that the data sent follows the defined schema:

json

```
{  
  "DeviceID": "device-123",  
  "Timestamp": "2024-09-19T12:34:56Z",  
  "SignalStrength": 85,  
  "CallDropRate": 2.3,  
  "DataTransferSpeed": 56  
}
```

3. Secure the Event Hub:

- **Access Policies:**

- Create **Shared Access Policies** with appropriate permissions (e.g., **Send** for devices, **Listen** for consumers).

2. Stream Processing and Enrichment with Azure Databricks

Objective: Process, clean, and enrich streaming data in real-time.

Steps:

1. Set Up Azure Databricks Workspace:

- **Create a Databricks Workspace:**

- In the Azure Portal, search for **Azure Databricks** and create a new workspace.

- **Create a Cluster:**
 - Launch the Databricks workspace and create a **Cluster** with appropriate configurations (e.g., **Standard_DS3_v2** nodes).

2. Mount ADLS Gen2 in Databricks:

- **Mount ADLS Gen2:**

Replace placeholders with your ADLS Gen2 details

```
configs = {
    "fs.azure.account.auth.type": "OAuth",
    "fs.azure.account.oauth.provider.type":
"org.apache.hadoop.fs.azurebfs.oauth2.ClientCredsTokenProvider",
    "fs.azure.account.oauth2.client.id": "<YOUR_CLIENT_ID>",
    "fs.azure.account.oauth2.client.secret": "<YOUR_CLIENT_SECRET>",
    "fs.azure.account.oauth2.client.endpoint":
"https://login.microsoftonline.com/<YOUR_TENANT_ID>/oauth2/token"
}

dbutils.fs.mount(
    source = "abfss://transformed-
data@<YOUR_STORAGE_ACCOUNT>.dfs.core.windows.net/",
    mount_point = "/mnt/transformed-data",
    extra_configs = configs
)
```

3. Create a Streaming Job in Databricks:

- **Read from Event Hubs:**

```
from pyspark.sql import SparkSession

from pyspark.sql.functions import from_json, col

from pyspark.sql.types import StructType, StructField, StringType, TimestampType,
IntegerType, FloatType
```

```
# Define the schema
```

```
schema = StructType([  
    StructField("DeviceID", StringType(), True),  
    StructField("Timestamp", TimestampType(), True),  
    StructField("SignalStrength", IntegerType(), True),  
    StructField("CallDropRate", FloatType(), True),  
    StructField("DataTransferSpeed", IntegerType(), True)  
])
```

```
# Read from Event Hubs
```

```
event_hub_connection_string = "<YOUR_EVENT_HUB_CONNECTION_STRING>"
```

```
df = (  
    spark.readStream  
    .format("eventhubs")  
    .option("eventhubs.connectionString", event_hub_connection_string)  
    .load()  
)
```

```
# Convert binary data to string
```

```
df = df.selectExpr("CAST(body AS STRING) as json_str")
```

```
# Parse JSON
```

```
df = df.select(from_json(col("json_str"), schema).alias("data")).select("data.*")
```

```
# Data Cleansing: Filter out invalid records
```

```
df_clean = df.filter((col("SignalStrength").isNotNull()) &  
    (col("CallDropRate").isNotNull()) &
```

```
(col("DataTransferSpeed").isNotNull()))
```

- **Enrich the Data (Optional):**

- Join with reference data, add calculated fields, etc.

Example: Add a calculated field

```
df_enriched = df_clean.withColumn("PerformanceScore",  
    (col("SignalStrength") * 0.5) +  
    ((1 - col("CallDropRate")) * 0.3) +  
    (col("DataTransferSpeed") * 0.2))
```

- **Write to ADLS Gen2 in Delta Format:**

```
query = (  
    df_enriched.writeStream  
    .format("delta")  
    .option("checkpointLocation", "/mnt/transformed-  
data/checkpoints/network_performance")  
    .option("path", "/mnt/transformed-data/network_performance")  
    .outputMode("append")  
    .start()  
)
```

```
query.awaitTermination()
```

3. Storage in Azure Data Lake Storage (ADLS) Gen2

Objective: Persist processed and enriched data for further analysis and reporting.

Steps:

1. Data Storage:

- The streaming job writes data to ADLS Gen2 in **Delta Lake** format, enabling ACID transactions and scalable storage.
- Data is stored under the path: /mnt/transformed-data/network_performance.

2. Access Control:

- Ensure that appropriate **Access Control Lists (ACLs)** are set on the ADLS Gen2 containers to secure data access.

4. Real-time Alerting for Network Anomalies

Objective: Detect anomalies in network performance metrics and trigger alerts.

Steps:

1. Define Anomaly Detection Logic in Databricks:

- Implement logic to identify anomalies based on predefined thresholds or statistical methods.

```
from pyspark.sql.functions import when
```

```
# Define thresholds
```

```
SIGNAL_STRENGTH_THRESHOLD = 50
```

```
CALL_DROP_RATE_THRESHOLD = 3.0
```

```
DATA_TRANSFER_SPEED_THRESHOLD = 40
```

```
# Identify anomalies
```

```
df_anomalies = df_enriched.filter(  
    (col("SignalStrength") < SIGNAL_STRENGTH_THRESHOLD) |  
    (col("CallDropRate") > CALL_DROP_RATE_THRESHOLD) |  
    (col("DataTransferSpeed") < DATA_TRANSFER_SPEED_THRESHOLD)
```

)

2. Write Anomalies to a Separate Delta Table:

```
anomaly_query = (  
    df_anomalies.writeStream  
  
    .format("delta")  
  
    .option("checkpointLocation", "/mnt/transformed-  
data/checkpoints/network_anomalies")  
  
    .option("path", "/mnt/transformed-data/network_anomalies")  
  
    .outputMode("append")  
  
    .start()  
)
```

```
anomaly_query.awaitTermination()
```

3. Trigger Alerts Using Azure Functions:

- **Create an Azure Function:**
 - Develop an Azure Function that listens to the network_anomalies Delta table or is triggered by events from Event Hubs.
- **Example: Azure Function to Send Alerts via Email or Teams:**

```
import requests
```

```
import json
```

```
import os
```

```
from azure.storage.blob import BlobServiceClient
```

```
def main(req: func.HttpRequest) -> func.HttpResponse:
```

```

device_id = req.params.get('DeviceID')
timestamp = req.params.get('Timestamp')
signal_strength = req.params.get('SignalStrength')
call_drop_rate = req.params.get('CallDropRate')
data_transfer_speed = req.params.get('DataTransferSpeed')

alert_message = f"Anomaly detected for Device {device_id} at {timestamp}.\n" \
    f"Signal Strength: {signal_strength}\n" \
    f"Call Drop Rate: {call_drop_rate}\n" \
    f>Data Transfer Speed: {data_transfer_speed}"

# Example: Send alert to Microsoft Teams via Incoming Webhook
teams_webhook_url = os.getenv("TEAMS_WEBHOOK_URL")
headers = {'Content-Type': 'application/json'}
payload = {"text": alert_message}
response = requests.post(teams_webhook_url, headers=headers,
data=json.dumps(payload))

if response.status_code == 200:
    return func.HttpResponse("Alert sent successfully.", status_code=200)
else:
    return func.HttpResponse(f"Failed to send alert: {response.text}", status_code=500)

```

- **Integrate with Databricks:**
 - Modify the Databricks streaming job to call the Azure Function when an anomaly is detected.

```
import requests
```



```

def send_alert(device_id, timestamp, signal_strength, call_drop_rate,
data_transfer_speed):

    url = "https://<YOUR_FUNCTION_APP>.azurewebsites.net/api/SendAlert"

    params = {

        "DeviceID": device_id,

        "Timestamp": timestamp,

        "SignalStrength": signal_strength,

        "CallDropRate": call_drop_rate,

        "DataTransferSpeed": data_transfer_speed

    }

    response = requests.get(url, params=params)

    if response.status_code != 200:

        print(f"Failed to send alert: {response.text}")


# Apply the function to each anomaly record

from pyspark.sql.functions import udf

from pyspark.sql.types import StringType


def trigger_alert(device_id, timestamp, signal_strength, call_drop_rate,
data_transfer_speed):

    send_alert(device_id, timestamp, signal_strength, call_drop_rate,
data_transfer_speed)


trigger_alert_udf = udf(trigger_alert, StringType())


df_anomalies.foreach(lambda row: trigger_alert(row.DeviceID, row.Timestamp,
row.SignalStrength, row.CallDropRate, row.DataTransferSpeed))

```

5. Visualization with Power BI/Fabric

Objective: Create real-time dashboards to monitor network performance metrics.

Steps:

1. Connect Power BI to ADLS Gen2 or Azure Databricks:

○ **Option 1: Connect Directly to Databricks:**

- In Power BI Desktop, select **Azure Databricks** as a data source.
- Provide the Databricks workspace URL and authentication token.
- Import the Delta tables (e.g., network_performance).

○ **Option 2: Connect via ADLS Gen2:**

- In Power BI Desktop, use the **Azure Data Lake Storage Gen2** connector.
- Provide the storage account details and access credentials.
- Import the Delta tables.

2. Create Real-time Dashboards:

○ **Design Visualizations:**

- **Signal Strength Trend:** Line chart showing signal strength over time.
- **Call Drop Rate Analysis:** Bar chart displaying call drop rates per device.
- **Data Transfer Speed Comparison:** Grouped bar chart comparing data transfer speeds across devices or time periods.
- **Geographical Heatmap:** Heatmap showing signal strength variations across different regions.
- **Real-time Alerts:** KPI cards or alert indicators showing current anomalies.

○ **Set Up Streaming Datasets (if using Direct Streaming):**

- Use **Push Datasets** or **Streaming Tiles** in Power BI to enable real-time updates.
- Configure **Power BI** to refresh data at short intervals (e.g., every minute).

3. Publish and Share Dashboards:

- Publish the Power BI reports to the **Power BI Service**.
 - Share dashboards with stakeholders and set up **Row-Level Security (RLS)** if needed.
-

6. Monitoring and Automation

Objective: Ensure the streaming pipeline runs smoothly and handle any failures or performance issues.

Steps:

1. Monitor Streaming Jobs in Databricks:

- **Databricks Workspace:**
 - Navigate to **Jobs** and monitor the status of your streaming jobs.
 - Check **Streaming Queries** under **Clusters** for real-time monitoring.

2. Set Up Alerts with Azure Monitor:

- **Create Metrics and Alerts:**
 - Use **Azure Monitor** to track metrics such as Event Hub throughput, Databricks job status, and storage utilization.
 - Configure alerts to notify you of any issues (e.g., job failures, high latency).

3. Automate Recovery and Scaling:

- **Auto-scaling in Databricks:**
 - Configure **Auto-scaling** for your Databricks clusters to handle varying workloads.
- **Retry Logic:**
 - Implement retry mechanisms in your streaming jobs and alerting functions to handle transient failures.

4. Use Azure Data Factory for Orchestration (Optional):

- While **Azure Databricks** handles the stream processing, you can use **Azure Data Factory (ADF)** to orchestrate additional workflows or manage dependencies between batch and streaming pipelines.
-

Schema Definition

Network Performance Schema:

Field	Data Type	Description
DeviceID	String	Unique identifier for the network device
Timestamp	Timestamp	Time when the data was recorded
SignalStrength	Integer	Signal strength measured in dBm or a similar unit
CallDropRate	Float	Percentage rate at which calls are dropped
DataTransferSpeed	Integer	Data transfer speed measured in Mbps or similar
PerformanceScore	Float (Optional)	Calculated performance score for anomaly detection

Code Examples

1. Streaming ETL Job in Databricks

```
from pyspark.sql import SparkSession

from pyspark.sql.functions import from_json, col, to_timestamp

from pyspark.sql.types import StructType, StructField, StringType, TimestampType, IntegerType, FloatType

# Initialize Spark Session

spark =
SparkSession.builder.appName("NetworkPerformanceStreaming").getOrCreate()

# Define schema

schema = StructType([

    StructField("DeviceID", StringType(), True),

    StructField("Timestamp", TimestampType(), True),
```

```
    StructField("SignalStrength", IntegerType(), True),  
    StructField("CallDropRate", FloatType(), True),  
    StructField("DataTransferSpeed", IntegerType(), True)  
  ])
```

```
# Read from Event Hubs
```

```
event_hub_connection_string =  
"Endpoint=sb://<YOUR_EVENT_HUB_NAMESPACE>.servicebus.windows.net;/SharedA  
ccessKeyName=<KEY_NAME>;SharedAccessKey=<KEY_VALUE>;EntityPath=network-  
performance-hub"
```

```
df = (  
    spark.readStream  
        .format("eventhubs")  
        .option("eventhubs.connectionString", event_hub_connection_string)  
        .load()  
)
```

```
# Convert binary data to string
```

```
df = df.selectExpr("CAST(body AS STRING) as json_str")
```

```
# Parse JSON
```

```
df = df.select(from_json(col("json_str"), schema).alias("data")).select("data.*")
```

```
# Data Cleansing
```

```
df_clean = df.filter((col("SignalStrength").isNotNull()) &  
    (col("CallDropRate").isNotNull()) &  
    (col("DataTransferSpeed").isNotNull()))
```

```
# Enrich Data
```

```
df_enriched = df_clean.withColumn("PerformanceScore",  
    (col("SignalStrength") * 0.5) +  
    ((1 - col("CallDropRate")) * 0.3) +  
    (col("DataTransferSpeed") * 0.2))
```

```
# Write to Delta Lake
```

```
query = (  
    df_enriched.writeStream  
        .format("delta")  
        .option("checkpointLocation", "/mnt/transformed-  
data/checkpoints/network_performance")  
        .option("path", "/mnt/transformed-data/network_performance")  
        .outputMode("append")  
        .start()  
)
```

```
query.awaitTermination()
```

2. Azure Function for Alerting

Function Code ():

```
import logging
```

```
import requests
```

```
import os
```

```
import azure.functions as func
```

```
def main(req: func.HttpRequest) -> func.HttpResponse:
```

```
logging.info('Azure Function triggered for network anomaly alert.')
```

```
try:
```

```
    # Extract query parameters
```

```
    device_id = req.params.get('DeviceID')
```

```
    timestamp = req.params.get('Timestamp')
```

```
    signal_strength = req.params.get('SignalStrength')
```

```
    call_drop_rate = req.params.get('CallDropRate')
```

```
    data_transfer_speed = req.params.get('DataTransferSpeed')
```

```
    if not all([device_id, timestamp, signal_strength, call_drop_rate,
data_transfer_speed]):
```

```
        return func.HttpResponse("Missing parameters", status_code=400)
```

```
    # Create alert message
```

```
    alert_message = f"""
```

```
    **Network Anomaly Detected**
```

```
    **Device ID**: {device_id}
```

```
    **Timestamp**: {timestamp}
```

```
    **Signal Strength**: {signal_strength}
```

```
    **Call Drop Rate**: {call_drop_rate}%
```

```
    **Data Transfer Speed**: {data_transfer_speed} Mbps
```

```
    Please investigate the issue immediately.
```

```
    """
```

```
    # Send alert to Microsoft Teams via Incoming Webhook
```

```

teams_webhook_url = os.getenv("TEAMS_WEBHOOK_URL")

headers = {'Content-Type': 'application/json'}

payload = {
    "text": alert_message
}

response = requests.post(teams_webhook_url, headers=headers, json=payload)

if response.status_code == 200:
    return func.HttpResponse("Alert sent successfully.", status_code=200)
else:
    logging.error(f"Failed to send alert: {response.text}")
    return func.HttpResponse(f"Failed to send alert: {response.text}",
status_code=500)

except Exception as e:
    logging.error(f"Error in alert function: {str(e)}")
    return func.HttpResponse(f"Error: {str(e)}", status_code=500)

```

Function Configuration:

- **Environment Variables:**
 - TEAMS_WEBHOOK_URL: URL of the Microsoft Teams Incoming Webhook connector.
- **HTTP Trigger:**
 - Configure the function to be triggered via HTTP requests.

Triggering the Function from Databricks:

```
import requests
```



```
def send_alert(device_id, timestamp, signal_strength, call_drop_rate,
data_transfer_speed):

    url = "https://<YOUR_FUNCTION_APP>.azurewebsites.net/api/SendAlert"

    params = {

        "DeviceID": device_id,

        "Timestamp": timestamp,

        "SignalStrength": signal_strength,

        "CallDropRate": call_drop_rate,

        "DataTransferSpeed": data_transfer_speed

    }

    response = requests.get(url, params=params)

    if response.status_code != 200:

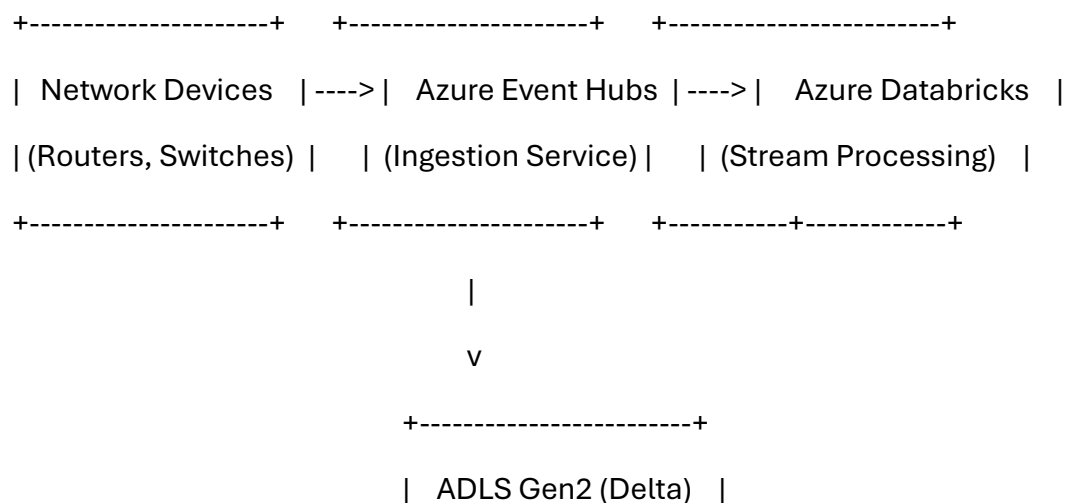
        print(f"Failed to send alert: {response.text}")

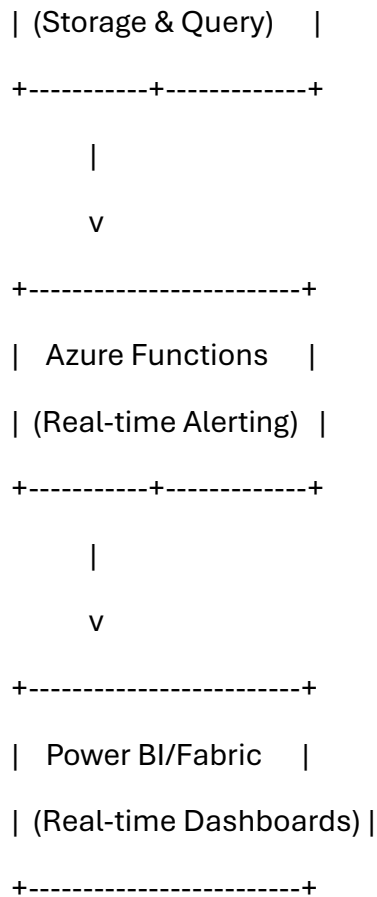
# Apply the function to each anomaly record

df_anomalies.foreach(lambda row: send_alert(row.DeviceID, row.Timestamp,
row.SignalStrength, row.CallDropRate, row.DataTransferSpeed))
```

Diagram of the Streaming Pipeline

plaintext





Detailed Configuration Steps

1. Setting Up Azure Event Hubs

1. Create an Event Hubs Namespace:

- **Navigate to Azure Portal > Create a Resource > Event Hubs.**
- **Fill in the Details:**
 - **Name:** telecom-event-hubs
 - **Pricing Tier:** Choose based on throughput needs (e.g., **Standard**).
 - **Resource Group:** Select or create a new one.
- **Review and Create.**

2. Create an Event Hub:

- **Within the Namespace,** select **Event Hubs > + Event Hub.**
- **Name:** network-performance-hub.
- **Configure Partitions:** Set based on expected load (e.g., 4 partitions).

- **Capture:** Optionally enable to automatically capture streaming data to ADLS Gen2.

3. Obtain Connection String:

- **Shared Access Policies:**
 - Under **Settings > Shared Access Policies**, select **RootManageSharedAccessKey** or create a new policy with **Send** permissions.
- **Connection String:** Copy the **Primary Connection String** for use in Databricks.

2. Configuring Azure Databricks for Streaming

1. Create and Configure the Cluster:

- **Databricks Workspace > Clusters > Create Cluster.**
- **Cluster Name:** network-performance-cluster.
- **Spark Version:** Choose the latest stable version (e.g., 7.3.x).
- **Node Type:** Standard_DS3_v2.
- **Workers:** 2 (adjust based on data volume).
- **Libraries:** Install necessary libraries (e.g., azure-eventhubs-spark).

2. Mount ADLS Gen2:

- **Notebook Cell:**

```
configs = {
    "fs.azure.account.auth.type": "OAuth",
    "fs.azure.account.oauth.provider.type":
"org.apache.hadoop.fs.azurebfs.oauth2.ClientCredsTokenProvider",
    "fs.azure.account.oauth2.client.id": "<YOUR_CLIENT_ID>",
    "fs.azure.account.oauth2.client.secret": "<YOUR_CLIENT_SECRET>",
    "fs.azure.account.oauth2.client.endpoint":
"https://login.microsoftonline.com/<YOUR_TENANT_ID>/oauth2/token"
}
```

```

dbutils.fs.mount(
    source = "abfss://transformed-
data@<YOUR_STORAGE_ACCOUNT>.dfs.core.windows.net/",
    mount_point = "/mnt/transformed-data",
    extra_configs = configs
)

```

3. Develop the Streaming Notebook:

- **Streaming Logic:**

```

from pyspark.sql import SparkSession

from pyspark.sql.functions import from_json, col, to_timestamp

from pyspark.sql.types import StructType, StructField, StringType, TimestampType,
IntegerType, FloatType

# Define the schema

schema = StructType([
    StructField("DeviceID", StringType(), True),
    StructField("Timestamp", TimestampType(), True),
    StructField("SignalStrength", IntegerType(), True),
    StructField("CallDropRate", FloatType(), True),
    StructField("DataTransferSpeed", IntegerType(), True)
])

# Read from Event Hubs

event_hub_connection_string = "Endpoint=sb://telecom-event-
hubs.servicebus.windows.net/;SharedAccessKeyName=<KEY_NAME>;SharedAccessK
ey=<KEY_VALUE>;EntityPath=network-performance-hub"

```

```
df = (  
    spark.readStream  
    .format("eventhubs")  
    .option("eventhubs.connectionString", event_hub_connection_string)  
    .load()  
)
```

Convert binary data to string

```
df = df.selectExpr("CAST(body AS STRING) as json_str")
```

Parse JSON

```
df = df.select(from_json(col("json_str"), schema).alias("data")).select("data.*")
```

Data Cleansing

```
df_clean = df.filter((col("SignalStrength").isNotNull()) &  
    (col("CallDropRate").isNotNull()) &  
    (col("DataTransferSpeed").isNotNull()))
```

Enrich Data

```
df_enriched = df_clean.withColumn("PerformanceScore",  
    (col("SignalStrength") * 0.5) +  
    ((1 - col("CallDropRate")) * 0.3) +  
    (col("DataTransferSpeed") * 0.2))
```

Write to Delta Lake

```
query = (  
    df_enriched.writeStream
```

```
.format("delta")

.option("checkpointLocation", "/mnt/transformed-
data/checkpoints/network_performance")

.option("path", "/mnt/transformed-data/network_performance")

.outputMode("append")

.start()

)
```

query.awaitTermination()

4. Run the Streaming Job:

- Execute the notebook to start the streaming job. Ensure that the cluster remains active to continue processing incoming data.

3. Real-time Alerting with Azure Functions

Objective: Automatically trigger alerts when anomalies are detected in network performance metrics.

Steps:

1. Create an Azure Function App:

- **Navigate to Azure Portal > Create a Resource > Compute > Function App.**
- **Fill in the Details:**
 - **Name:** NetworkAnomalyAlertFunction.
 - **Runtime Stack:** .
 - **Hosting:** Choose an appropriate plan (e.g., **Consumption Plan** for serverless).
 - **Storage Account:** Create or use an existing one.
 - **Region:** Choose the same region as other resources for lower latency.
- **Review and Create.**

2. Develop the Function:

- **Function Code ():**

```
import logging

import requests

import os

import azure.functions as func

def main(req: func.HttpRequest) -> func.HttpResponse:

    logging.info('Azure Function triggered for network anomaly alert.')

    try:

        # Extract query parameters

        device_id = req.params.get('DeviceID')

        timestamp = req.params.get('Timestamp')

        signal_strength = req.params.get('SignalStrength')

        call_drop_rate = req.params.get('CallDropRate')

        data_transfer_speed = req.params.get('DataTransferSpeed')

        if not all([device_id, timestamp, signal_strength, call_drop_rate,
data_transfer_speed]):

            return func.HttpResponse("Missing parameters", status_code=400)

        # Create alert message

        alert_message = f""

        **Network Anomaly Detected**

        **Device ID**: {device_id}
```

```
**Timestamp**: {timestamp}

**Signal Strength**: {signal_strength}

**Call Drop Rate**: {call_drop_rate}%

**Data Transfer Speed**: {data_transfer_speed} Mbps
```

Please investigate the issue immediately.

```
"""
```

```
# Send alert to Microsoft Teams via Incoming Webhook
teams_webhook_url = os.getenv("TEAMS_WEBHOOK_URL")
headers = {'Content-Type': 'application/json'}
payload = {
    "text": alert_message
}

response = requests.post(teams_webhook_url, headers=headers, json=payload)

if response.status_code == 200:
    return func.HttpResponse("Alert sent successfully.", status_code=200)
else:
    logging.error(f"Failed to send alert: {response.text}")
    return func.HttpResponse(f"Failed to send alert: {response.text}",
status_code=500)

except Exception as e:
    logging.error(f"Error in alert function: {str(e)}")
    return func.HttpResponse(f"Error: {str(e)}", status_code=500)
```

- **Configure Environment Variables:**

- In the Function App settings, under **Configuration**, add a new application setting:

- **Name:** TEAMS_WEBHOOK_URL
- **Value:** Your Microsoft Teams Incoming Webhook URL.

3. Integrate Azure Databricks with Azure Functions:

- **Modify the Streaming Job to Call the Azure Function:**

```
import requests
```

```
def send_alert(device_id, timestamp, signal_strength, call_drop_rate,
data_transfer_speed):
```

```
    url =
    "https://<YOUR_FUNCTION_APP>.azurewebsites.net/api/NetworkAnomalyAlertFunction"
```

```
    params = {
        "DeviceID": device_id,
        "Timestamp": timestamp,
        "SignalStrength": signal_strength,
        "CallDropRate": call_drop_rate,
        "DataTransferSpeed": data_transfer_speed
    }
```

```
    response = requests.get(url, params=params)
```

```
    if response.status_code != 200:
```

```
        print(f"Failed to send alert: {response.text}")
```

```
# Apply the function to each anomaly record
```

```
df_anomalies = df_enriched.filter(
```

```
    (col("SignalStrength") < SIGNAL_STRENGTH_THRESHOLD) |
```

```
    (col("CallDropRate") > CALL_DROP_RATE_THRESHOLD) |
```

```
    (col("DataTransferSpeed") < DATA_TRANSFER_SPEED_THRESHOLD)
```

)

```
query_anomalies = (  
    df_anomalies.writeStream  
        .foreach(lambda row: send_alert(row.DeviceID, row.Timestamp, row.SignalStrength,  
row.CallDropRate, row.DataTransferSpeed))  
        .outputMode("append")  
        .start()  
)
```

```
query_anomalies.awaitTermination()
```

4. Visualization with Power BI/Fabric

Objective: Create real-time dashboards to monitor network performance and visualize anomalies.

Steps:

1. **Connect Power BI to ADLS Gen2 or Databricks:**

- **Option 1: Connect Directly to Databricks:**
 - **In Power BI Desktop:**
 - Click on **Get Data > Azure > Azure Databricks**.
 - Enter the **Databricks workspace URL** and **Personal Access Token**.
 - Select the Delta tables (e.g., network_performance).
- **Option 2: Connect via ADLS Gen2:**
 - **In Power BI Desktop:**
 - Click on **Get Data > Azure > Azure Data Lake Storage Gen2**.
 - Enter the **ADLS Gen2 URL** and authenticate using OAuth or Access Key.

- Navigate to the path /transformed-data/network_performance and load the Delta tables.

2. Create Real-time Dashboards:

- **Design Visualizations:**

- **Signal Strength Trend:**

- **Visualization:** Line Chart.
 - **X-Axis:** Timestamp.
 - **Y-Axis:** SignalStrength.
 - **Filter:** By DeviceID or aggregate across devices.

- **Call Drop Rate Analysis:**

- **Visualization:** Bar Chart.
 - **X-Axis:** DeviceID.
 - **Y-Axis:** CallDropRate.

- **Data Transfer Speed Comparison:**

- **Visualization:** Grouped Bar Chart.
 - **X-Axis:** DeviceID.
 - **Y-Axis:** DataTransferSpeed.

- **Geographical Heatmap:**

- **Visualization:** Filled Map or Shape Map.
 - **Location Data:** If available (requires device location data).

- **Real-time Alerts:**

- **Visualization:** KPI Cards or Conditional Formatting Tables.
 - **Indicators:** Display latest anomalies or counts of anomalies.

- **Performance Score Overview:**

- **Visualization:** Scatter Plot.
 - **X-Axis:** CallDropRate.
 - **Y-Axis:** PerformanceScore.
 - **Color:** Indicate severity or device category.

3. Enable Real-time Data Streaming:

- **Using Power BI Streaming Datasets:**
 - Set up a **Streaming Dataset** in Power BI for real-time data.
 - Use **Power BI REST API** or **Azure Stream Analytics** to push real-time data to Power BI.
- **Configure Automatic Refresh:**
 - Schedule **DirectQuery** or **Live Connection** for near real-time updates.
 - Ensure that data sources are optimized for frequent querying.

4. Publish and Share Dashboards:

- **Publish:**
 - Publish the Power BI report to the **Power BI Service**.
- **Share:**
 - Share dashboards with stakeholders and set up **Row-Level Security (RLS)** if needed.
- **Mobile Access:**
 - Enable mobile access to dashboards for on-the-go monitoring.

5. Monitoring and Automation

Objective: Ensure the streaming pipeline operates smoothly and efficiently.

Steps:

1. Monitor Streaming Jobs in Databricks:

- **Databricks Workspace:**
 - Navigate to **Jobs** and monitor the status of your streaming jobs.
 - Use **Streaming Query** monitoring to check for any backlogs or issues.

2. Set Up Alerts with Azure Monitor:

- **Create Metrics and Alerts:**
 - Monitor **Event Hubs** metrics (e.g., incoming requests, errors).

- Monitor **Databricks** job metrics (e.g., processing rates, errors).
- Configure alerts for any anomalies or failures in the pipeline.

3. Use Azure Data Factory (Optional):

- **Orchestrate Additional Workflows:**
 - Use **ADF** to manage dependencies between batch and streaming pipelines.
 - Schedule data ingestion tasks or integrate with other Azure services as needed.

4. Implement Auto-scaling:

- **Databricks Auto-scaling:**
 - Configure your Databricks clusters to auto-scale based on workload to handle variable data volumes.

5. Optimize Performance:

- **Efficient Data Partitioning:**
 - Ensure Delta tables are properly partitioned for optimized query performance.
- **Caching:**
 - Use caching strategies in Databricks for frequently accessed data.

Conclusion

By following this end-to-end solution, you can effectively implement **Scenario 2: Real-time Network Monitoring** within your telecom data engineering project using Azure Cloud services. This setup ensures that network performance metrics are ingested, processed, stored, and visualized in real-time, providing immediate insights and enabling proactive responses to network anomalies.

Key Benefits:

- **Real-time Insights:** Immediate visibility into network performance allows for swift action on issues.
- **Scalability:** Azure services scale seamlessly to handle varying data volumes and processing needs.

- **Automation:** Automated data pipelines reduce manual intervention and enhance reliability.
- **Integration:** Seamless integration with Azure's ecosystem ensures robust security, monitoring, and management.
- **Visualization:** Power BI/Fabric provides intuitive dashboards for stakeholders to monitor network health continuously.